

PROJECT BLUEPRINT — ML APPLICATION


Personalized Health Insights System

An AI-powered platform that learns from a user's daily health habits — sleep, activity, nutrition, stress — and delivers personalized, evidence-backed recommendations to improve wellbeing over time.

WHAT IS THIS PROJECT?

- 🎯 Personalized health coaching via ML
- 📊 Multi-source habit tracking
- 🧠 Predictive risk modeling
- 💬 AI-generated daily insights
- 📱 Web + Mobile interface

HOW IT WORKS — DATA PIPELINE



WHY BUILD THIS?

Preventive healthcare gap

Most people only see a doctor reactively. This system acts as a proactive, always-on health advisor.

Personalization at scale

Generic health advice is ignored. ML enables truly individual recommendations based on each user's unique patterns.

Wearable data is underused

Millions of users generate rich biometric data daily; most apps display it without making it actionable.

Behavior change science

ML can model habit loops, detect declining trends early, and suggest timely nudges to improve outcomes.

PROJECT MODULES

1 User Profile & Onboarding

Collects demographic data, health goals, medical history, and wearable device connections. Establishes the user's baseline health fingerprint.

React FastAPI PostgreSQL

2 Data Ingestion & Integration

Pulls data from wearables (Fitbit, Apple Health, Garmin), manual logs, and food tracking APIs. Normalizes and stores in a time-series database.

Apache Kafka InfluxDB REST/OAuth2

3 Feature Engineering

4 ML Modeling Engine

Transforms raw data into ML-ready features: sleep debt calculations, weekly activity trends, caloric balance, HRV patterns, and stress indices.

Python Pandas Scikit-learn

Core intelligence layer. Multiple models: collaborative filtering for recommendations, LSTM for time-series trend prediction, and classification for risk flags.

TensorFlow XGBoost MLflow

5 Recommendation Engine

Translates model outputs into actionable insights: sleep tips, workout suggestions, dietary adjustments, hydration reminders — ranked by predicted impact.

LLM API Rule engine A/B testing

6 Dashboard & Visualization

User-facing interface showing trends, weekly health scores, progress toward goals, and daily AI-generated insight cards with explanation.

React Recharts Tailwind

7 Notification & Nudge System

Sends timely push/email alerts when trends decline or a health goal is at risk. Uses optimal timing prediction to maximize engagement.

Firebase FCM Celery Redis

8 Privacy, Security & Compliance

HIPAA-aware data handling, end-to-end encryption, anonymization pipelines, user data export/deletion, and audit logging for all model decisions.

AES-256 GDPR HIPAA

TECHNOLOGY STACK

Layer	Technologies	Purpose
Frontend	React 18 Next.js Recharts Tailwind	Responsive dashboard, charts, mobile PWA
Backend	FastAPI Python 3.11 Celery Redis	REST APIs, async task queues, caching
ML / AI	TensorFlow 2 Scikit-learn XGBoost MLflow Hugging Face	Model training, tracking, NLP for insights
Databases	PostgreSQL InfluxDB MongoDB	Relational, time-series, and document stores
Infrastructure	Docker Kubernetes AWS/GCP Kafka GitHub Actions	Containerization, streaming, CI/CD

REQUIRED DATASETS & SOURCES

MIMIC-IV Clinical Database

Large-scale clinical records including vitals, labs, medications. Used for training risk models and health correlation baselines.

physionet.org/content/mimiciv

NHANES (CDC)

National Health and Nutrition Examination Survey — nutrition, physical activity, sleep data for 50,000+ Americans. Ideal for population-level baselines.

cdc.gov/nchs/nhanes

Sleep-EDF (Physionet)

Annotated polysomnography recordings for sleep stage classification. Used to train sleep quality prediction models from wearable sensor data.

physionet.org/content/sleep-edf

Open Food Facts

Crowd-sourced food database with 2M+ products and full nutritional profiles. Powers the dietary logging and food recommendation modules.

world.openfoodfacts.org/data

PMData — Fitness Wearables Dataset

Real-world Fitbit + self-reported data from 16 participants over 5 months: steps, HR, sleep, mood, RPE. Best for wearable feature correlation experiments.

github.com/simula/pmdata

WESAD — Stress & Affect Detection

Multimodal physiological signals (chest + wrist) for stress/affect state classification. Trains the stress detection component of the ML engine.

uni-siegen.de/ubicomp/papers/ubi_jwe2018.pdf

SUGGESTED BUILD ROADMAP

- **Phase 1 — Foundation (Weeks 1–3)**
Set up project repo, database schema, user auth, and basic onboarding UI. Define data models and API contracts.
- **Phase 2 — Data Layer (Weeks 4–6)**
Build ingestion pipelines, wearable OAuth integrations, manual log forms, and data normalization scripts. Connect InfluxDB for time-series storage.
- **Phase 3 — ML Core (Weeks 7–11)**
Feature engineering, baseline models (sleep quality, activity classification), LSTM trend predictor, collaborative filtering for recommendations. Track with MLflow.
- **Phase 4 — Insights Layer (Weeks 12–14)**
Recommendation engine, LLM-generated insight cards, nudge logic, and A/B testing framework for recommendation variants.
- **Phase 5 — Dashboard & Polish (Weeks 15–17)**

Full React dashboard, charts, health score visualization, notification system, and progressive web app setup.

- Phase 6 — Security & Launch (Weeks 18–20)

Encryption, HIPAA audit, penetration testing, Kubernetes deployment, monitoring dashboards, and beta user onboarding.